

# COMS W4701: Artificial Intelligence

## Lecture 2: Intelligent Agents

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences


# Today

---

- Rational agents
- Task environments
- Agent programs


# Agents and Environments

- An **agent** perceives its environment through *sensors* and acts upon its environment through *actuators*
- Agents may be *controlled* or *autonomous*
- **Agent functions** map percepts to actions
- Other dependencies: Prior knowledge, past experience, goals, preferences, capabilities
- Environment may be arbitrarily general, but in practice limited to aspects that directly interact with the agent




# Examples of Agents

- Humans are agents!
  - Sensors: Vision, hearing, touch, smell, taste, proprioception
  - Actuators: Muscles, reflexes, changing brain state
- AI is more interested in “artificial” agents in complex environments that require nontrivial decision making
- Example: Vacuum cleaner (think Roomba)
- Environment: Square A and square B



# Example: Vacuum Cleaner World

- **Percepts:** Current square; is the square dirty?
- **Actions:** Move left, move right, clean, do nothing



- **Agent function:**

[A, IsClean]	Move right
[A, IsDirty]	Clean
[B, IsClean]	Move left
[B, IsDirty]	Clean
[[A, IsDirty], [A, IsClean]]	Move right
[[B, IsDirty], [A, IsClean]]	Move right
[[B, IsClean], [A, IsClean]]	Do nothing
.....	.....

# Rational Agents

---

- Environment state sequence can be evaluated by a **performance measure**
- Performance measures usually based on desired *outcomes*, not *behaviors*
- A **rational agent** selects an action to *maximize* its performance measure given percept sequence and in-built knowledge.
  - What performance measure makes our vacuum cleaner rational or irrational?
- Rational agents maximize *expected* performance, are not omniscient
- Rationality may involve info gathering, exploration, learning

# Task Environments

---

- A rational agent is a *solution* to a **task environment** problem
- **PEAS**: Performance measure, environment, actuators, sensors
- Vacuum cleaner task environment
  - P: Cleanliness, power usage, time taken
  - E: The small grid world
  - A: Wheels to move, filter to clean
  - S: “GPS”, cleanliness sensor

# Task Environment Properties

---

- Fully observable vs partially observable vs unobservable
  - Can agent sense all *relevant* information? Is internal state (memory) required?
- Single-agent vs multi-agent
  - Does maximization of performance measure depend on other agents' behaviors?
- Deterministic vs stochastic
  - Can we completely predict the next state of the environment?
- Episodic vs sequential
  - Do current decisions depend on past ones? Do current decisions affect future ones?
- Static vs dynamic
  - Does the environment change while the agent is thinking?
- Discrete vs continuous
  - Is number of states, actions, percepts, time, etc. finite?



# Examples of Environments

Environment	Partially / Fully Observable	Single- / Multi-Agent	Deterministic / Stochastic	Sequential / Episodic	Dynamic / Static	Continuous / Discrete
Vacuum cleaner world	Depends	Single	Deterministic	Sequential	Static	Discrete
Chess	Fully	Multi (adversarial)	Deterministic	Sequential	Static	Discrete
Self-driving car	Partially	Multi (cooperative)	Stochastic	Sequential	Dynamic	Continuous
Image classification	Fully	Single	Deterministic	Episodic	Static	Depends


# Agent Design

---

- Understanding the task environment tells us how to design our agents
- The more difficult the task environment, the more complex the agent
- Partially observable env -> agent requires memory / state
- Multi-agent env -> agent requires tracking of other agents
- Stochastic env -> agent must consider multiple scenarios or contingencies
- Sequential env -> agent must consider past and future states
- Dynamic env -> agent must maintain model of the world

# Agent Programs

- An **agent program** is an *implementation* of an agent function
- Specifies *how* something is computed rather than *what* needs to be computed
- A given agent program is not a universal solution!
- Program usefulness depends on hardware, tractability
- Agent programs also depend on the desired *solutions*
- E.g., *optimal* wrt some utility, approximately optimal, *satisficing* (“good enough”), *probable*

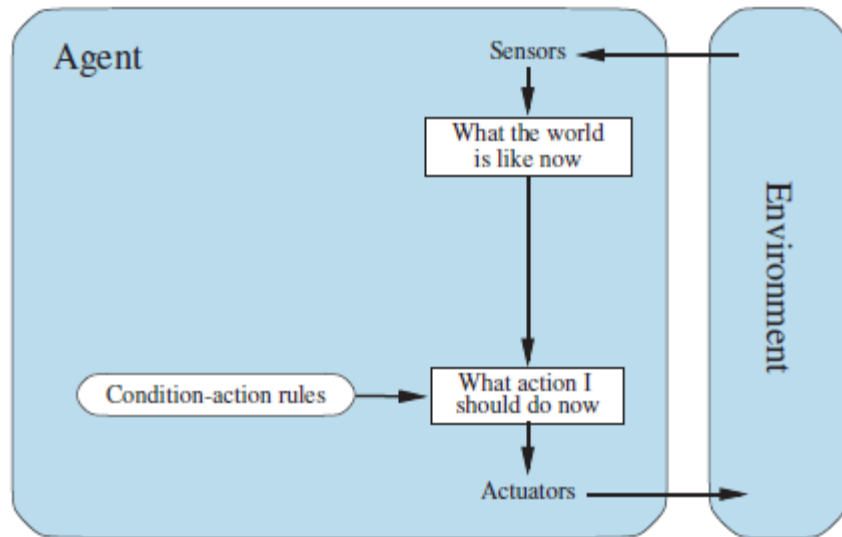


# Simple Reflex Agents

- Simple reflex agent: Use current percept only
- Percepts may be mapped to internal *states*
- Match state to condition-action (if-then) rules


```
function SIMPLE-REFLEX-AGENT(percept) returns an action
  persistent: rules, a set of condition-action rules

  state ← INTERPRET-INPUT(percept)
  rule ← RULE-MATCH(state, rules)
  action ← rule.ACTION
  return action
```



# Example: Vacuum Cleaner Robot

- Example: Vacuum cleaner robot as a reflex agent
- Use only the *current* (no past) percept
- What is its resultant behavior?



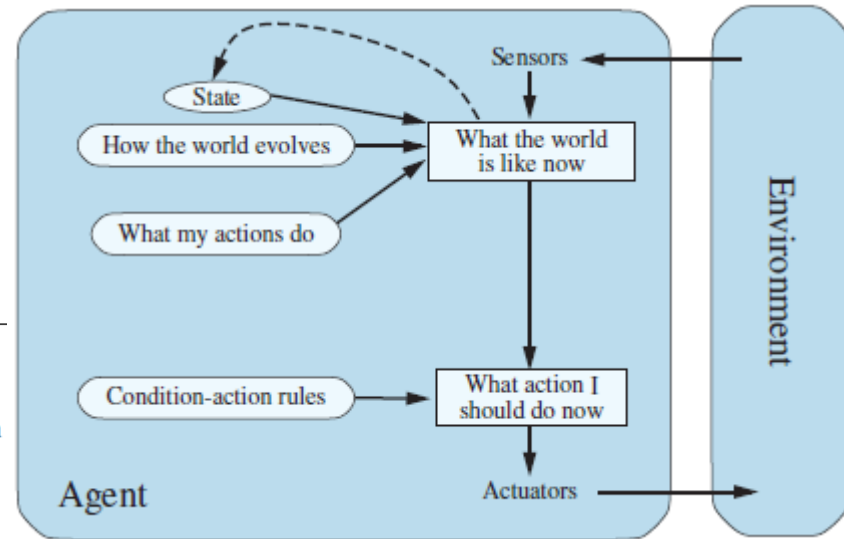
```
function Reflex-Vacuum-Agent([location, status]) returns action
  if status is dirty then return clean
  if location is A then return move right
  if location is B then return move left
  return do nothing
```

# Model-Based Reflex Agents

- What about partially observable environments?
- Maintain **internal state** and **transition model**
- May also have *sensor model* mapping percepts
- Use all information to *update* the state

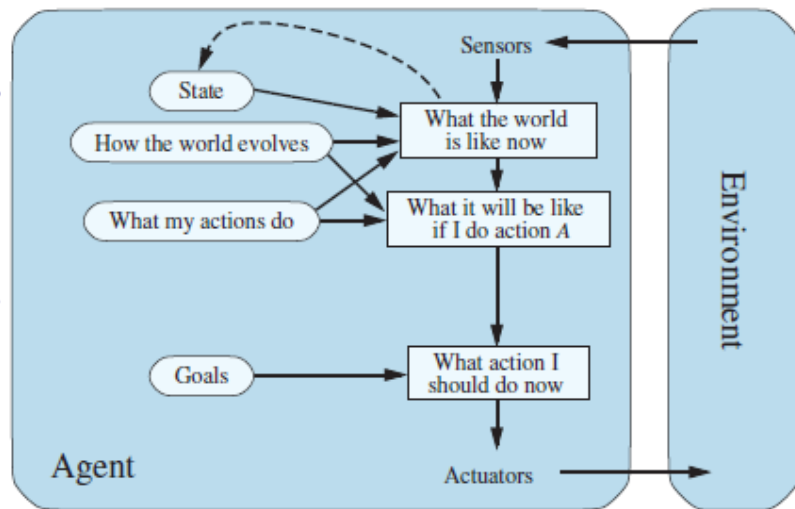
```
function MODEL-BASED-REFLEX-AGENT(percept) returns an action
  persistent: state, the agent's current conception of the world state
               model, a description of how the next state depends on current state and action
               rules, a set of condition-action rules
               action, the most recent action, initially none

  state ← UPDATE-STATE(state, action, percept, model)
  rule ← RULE-MATCH(state, rules)
  action ← rule.ACTION
  return action
```



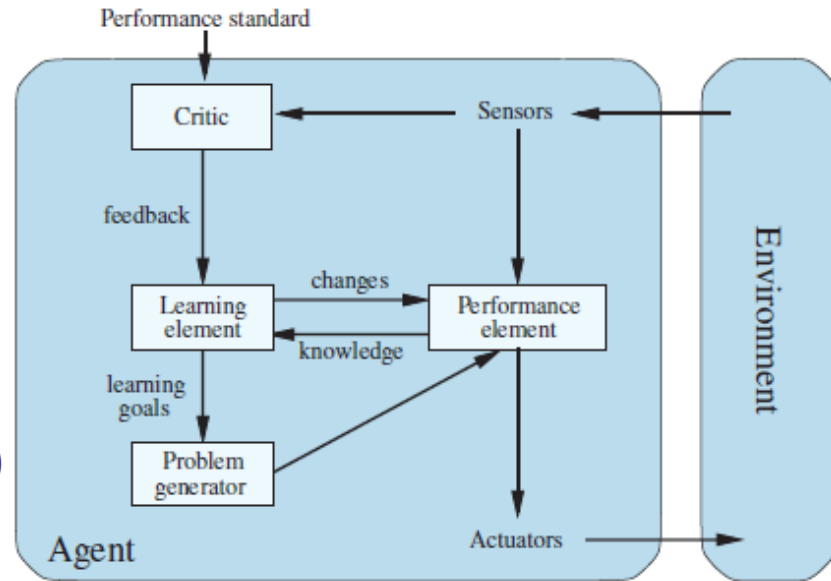
# Goal- and Utility-Based Agents

- Reflex agents are very rigid and predictable
- **Goal-based** agents try to achieve particular states
- Problems often solved using search and planning
- **Utility-based** agents can compare different states
- Utility functions map state to “desirability”
- Internalize the overall performance measure
- Utilities specify tradeoffs for competing goals
- Also useful in face of uncertainty



# Learning Agents

- **Learning agents** can be used to *create* or *improve* upon initial models in unknown envs
- A *learning element* retrieves knowledge from and then improves the *performance element*
- A *critic* evaluates the learning element according to a *performance standard* (measure)
- A *problem generator* suggests actions that can help gather new information and experiences





# Summary

---

- Agents interact with their environments through sensors and actuators
- Rational agents maximize expected performance measure
- PEAS descriptors define task environments and influence agent design
- Environment properties vary from easy to extremely challenging
- Agent programs may use current percept only, use a model, and/or try to achieve certain goals quantified by utilities
- Agent programs may also be created or improved via learning