

COMS W4701: Artificial Intelligence

Lecture 5: Constraint Satisfaction Problems

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences

Today

- Constraint satisfaction problems
- Backtracking search
- Constraint propagation and local consistency
- Heuristics on problem structure and constraints

States with Structure


- So far, we have reasoned explicitly in terms of states as “black boxes”
- Oftentimes, we can also describe states using a common set of *features*
- Features can be denoted using **variables** and assigned **values**
- We may have hard **constraints** or **preferences** over assignments
- The planning problem (search for goal) is now an *assignment* problem
- We can also apply more *general*, rather than problem-specific, heuristics

Constraint Satisfaction Problems

- Special structured search problems with 3 components
 - Variables (e.g., discrete, binary, Boolean, continuous, etc.): $X = \{X_1, \dots, X_n\}$
 - Domain for each variable (may be the same for some/all): $D = \{D_1, \dots, D_n\}$
 - Constraints over variables (e.g., unary, binary, trinary, etc.): $C = \{C_1, \dots, C_m\}$
- Constraints may be defined *implicitly*, e.g., using formulas, logic, or relations
- May also be defined *explicitly*, i.e., as a set of all allowed assignments
- Goal test: A **complete, consistent** assignment of values to each variable X_i from respective domain D_i s.t. all constraints C are satisfied


Example: Map Coloring

- Goal: Color a map so that no adjacent territories have the same color
- Variables: $X = \{WA, NT, Q, NSW, V, SA, T\}$
- Domains: $D_i = \{\text{red, green, blue}\}$
- Constraints: Implicit vs explicit representation
 - $C = \{WA \neq NT, WA \neq SA, NT \neq SA, NT \neq Q, SA \neq Q, Q \neq NSW, NSW \neq V\}$
 - $C = \{(WA, NT) \in \{(\text{red, green}), (\text{red, blue}), (\text{green, red}), (\text{green, blue}), \dots\}, \dots\}$



Example: n -Queens

- Place n queens on $n \times n$ board s.t. none share a row, column, or diagonal
- Can come up with more than one CSP representation
- CSP 1: A Boolean variable for each cell (n^2 total)
- Constraints: Every set of row, column, and diagonal cells contains at most one queen
- CSP 2: A variable X_i for row i with domain $\{0,1,2,3\}$ (4 total)
- Constraints: $\forall i, j, X_i \neq X_j$ and $X_j - X_i \neq |j - i|$



More Examples

- **Cryptarithmic**

- Variables: $\{T, W, O, F, U, R, C_1, C_2\}$

- Domains: $\{0, \dots, 9\}$

- Constraints: All variables different, and each column satisfies addition relation accounting for carryover variables C_1 and C_2

$$\begin{array}{r} T \ W \ O \\ + \ T \ W \ O \\ \hline F \ O \ U \ R \end{array}$$

- **Sudoku**

- Variables: One for each empty cell with domain $\{1, \dots, 9\}$

- Constraints: Different values in each row, col, 3x3 square

					8			4
	8	4		1	6			
			5			1		
1		3	8			9		
6		8				4		3
		2			9	5		1
		7			2			
			7	8		2	6	
2			3					


Backtracking Search


- Naïve idea: Simply test all complete assignments for consistency
- For n variables with domain size d , total number of assignments is d^n
- We don't have to check all of them, since some will violate constraints

- Better idea: *Incrementally* assign one variable at a time
- Only assign consistent values so that a constraint is never violated

- If we cannot progress further from a partial assignment, *backtrack* and change some current assignments to try alternatives

Backtracking Search

- Backtracking search is just a specific case of DFS!
 - States are partial assignments, goal is a complete assignment
 - Each search iteration tries to make a new assignment
 - If we cannot do so, DFS *backtracks* and considers the next closest partial assignment
 - Same time and space complexities as DFS
 - May be more efficient as constraints help *prune* invalid assignments from the search tree
- 



Constraint Propagation


- Before or during search, we can also *reduce the domains* of variables according to the constraints so that they satisfy **local consistency**
- **Node (domain) consistency:** Remove all domain values violating unary constraints
 - Should always be done prior to starting any search process
- **Arc consistency:** Remove all domain values violating binary constraints
 - X is consistent wrt Y if $\forall x \in D_1, \exists y \in D_2$ s.t. $(X = x, Y = y)$ is consistent
- To make X consistent with Y , check all values of and potentially remove values from D_1
- To make Y consistent with X , check all values of and potentially remove values from D_2

Example: Arc Consistency

- Variables A, B, C, D with domains $\{1,2,3,4\}$
- Constraints $c_1: A < B$ and $c_2: B < C$
- Make A consistent with B : $A = 4$ does not satisfy c_1 , so reduce A domain to $\{1,2,3\}$
- Make B consistent with A : $B = 1$ does not satisfy c_1 , so reduce B domain to $\{2,3,4\}$
- Make B consistent with C : $B = 4$ does not satisfy c_2 , so reduce B domain to $\{2,3\}$
- Make C consistent with B : $C = 1,2$ both cannot satisfy c_2 , so reduce C domain to $\{3,4\}$
- Since B 's domain changed, $A = 3$ no longer satisfies c_1 , so reduce A domain to $\{1,2\}$


AC-3 Algorithm

- CSP representation as a **constraint graph** where variable vertices are adjacent to constraint vertices
- Algorithm: Initialize a queue containing all constraints to check for consistency
- While queue not empty, pop a constraint and modify neighbor domains for consistency
- If domain of X changes, all constraint neighbors must be *re-added* to the queue to recheck




AC-3 with Search

- AC-3 complexity: Each constraint check takes $O(d^2)$ time
 - For each value in D_i , check each value in D_j (and vice versa)
 - There are m constraints, each may be checked up to d times
 - Overall time complexity is thus $O(md^3)$ —not too bad!
-
- AC-3 ends when entire CSP is arc-consistent or no assignment is possible
 - A arc-consistent CSP may have no solution discovered upon further search
-
- Can look at higher orders of consistency, but the tradeoff is more computation
 - **Path consistency:** $\forall (x_1, x_2) \in D_1 \times D_2, \exists x_3 \in D_3$ s.t. (x_1, x_2, x_3) is consistent



Variable Elimination

- Another class of strategies exploits and modifies the constraint graph *structure*
- We can *shrink* the graph by eliminating one or more variables
- We can replace a variable X with a new constraint specifying valid assignments for its neighbors
- Alternatively, we can assign X to different values in its domain, ensure arc consistency with neighbors, and then remove it from the graph
- In both cases, the resultant graph will have simpler structure than the original



Heuristics on Constraints

- **Bounds propagation:** If constraints are equality or inequality (e.g., *resource constraints*), we can use them to tighten finite domain bounds
 - Ex: X_1, X_2, X_3 , all with domains $\{1, 2, 3, 4, 5\}$
 - Constraint $\sum X_i = 13$: Reduce domains to $\{3, 4, 5\}$
 - Constraint $\sum X_i \leq 5$: Reduce domains to $\{1, 2, 3\}$
- If a CSP has multiple solutions due to *value symmetry* or assignment permutations, we can *introduce symmetry-breaking constraints* to help reduce domain sets
 - Ex: $X_1 \neq X_2$ can be replaced with $X_1 < X_2$ if it does not affect other variables/constraints
- **Constraint learning:** When we see inconsistent assignments in backtracking or constraint propagation, record it as a new constraint

Summary

- CSPs are structured search problems with variables, domains, constraints
- Solution consists of a complete, consistent assignment (goal, not path)
- Backtracking search performs incremental assignments while satisfying constraints, behaves like DFS
- Constraint propagation ensures local consistency of variable domains
- Other heuristics may involve modifying or adding new constraints