

COMS W4701: Artificial Intelligence

Lecture 7: Optimal Game Playing

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences

Today

- Adversarial search problems
- Minimax search
- Alpha-beta pruning
- Move ordering

Multi-Agent Environments

- Consider environment with multiple autonomous agents
- Each agent maintains own state about itself and other agents
- Each agent tries to maximize its own utility, dependent on env outcomes
- Fully **cooperative** if all agents share the same utility function
- Fully **competitive** if an agent effectively seeks to *minimize* others' utilities
 - **Zero-sum** if sum of all utilities in any outcome is always zero
- Most multi-agent environments (*games*) lie somewhere on this spectrum

Competitive Environments


- Several different approaches for competitive, multi-agent environments
- *Aggregate*, e.g., as in an economy; study overall effects and properties without worrying about individual agents
- Consider other agents to be part of a *nondeterministic* environment
- Our agent may perform better if we consider other agents' strategies
- Come up with a *strategy* using **adversarial search** and assuming that other agents are also rational actors

Computer Checkers, Chess, and Go


- 1994: *Chinook* declared computer champion in checkers against Tinsley
- 2007: Checkers is **solved** (predictable from any position assuming perfect play)
- 1997: *Deep Blue* defeats chess world champion Kasparov
- 2016: *AlphaGo* defeats go world champion Lee Sedol
- 2017: *AlphaZero* defeats *Stockfish*, AI chess champion
- 2018-today: *Stockfish* continues to improve using alpha-beta search methods
- *Leela Chess Zero* demonstrates comparable performance using methods based on *AlphaZero* (reinforcement learning, self-play)

Adversarial Search

- Assumption: *Perfect information* (fully observable) and turn-taking game
- Represent game progression (and later search) in a **game tree**
- Each layer (*ply*) of internal nodes and edges is controlled by one agent
- From a given state, the controlling agent may take one action (edge)
- Leaf nodes represent *terminal states*
- **Utility function** specifies utility values for each player at each terminal state



Game Trees



- Game trees can be *huge*, even for games like tic-tac-toe ($O(10^5)$ leaves)
- Chess has $O(10^{123})$, go has $O(10^{360})$!
- If a player wants to take optimal action, they can *search* the tree down to leaf nodes to see all possible utility outcomes


Minimax Values

- Suppose we have a 2-player, zero-sum game
- Utility function only assigns one value to a terminal state
- One player wants to maximize; other wants to minimize
- **Minimax value:** Value of a state assuming each player always plays *optimally* at every game state

$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

- We can compute all values by *searching* through the tree!

Minimax values:
computed *upward*
from terminal states



Terminal values:
part of the game

Minimax Search Algorithm

function MINIMAX-SEARCH(*game*, *state*) **returns** an action

$\text{player} \leftarrow \text{game.TO-MOVE}(\text{state})$

$\text{value}, \text{move} \leftarrow \text{MAX-VALUE}(\text{game}, \text{state})$

return *move*

Assuming root is MAX

function MAX-VALUE(*game*, *state*) **returns** a (*utility*, *move*) pair

if *game.IS-TERMINAL*(*state*) **then return** *game.UTILITY*(*state*, *player*), *null*

$v \leftarrow -\infty$

for each *a* **in** *game.ACTIONS*(*state*) **do**

$v2, a2 \leftarrow \text{MIN-VALUE}(\text{game}, \text{game.RESULT}(\text{state}, a))$

if $v2 > v$ **then**

$v, \text{move} \leftarrow v2, a$

return *v*, *move*

Compute and return max over
successors, all of which are MIN

function MIN-VALUE(*game*, *state*) **returns** a (*utility*, *move*) pair

if *game.IS-TERMINAL*(*state*) **then return** *game.UTILITY*(*state*, *player*), *null*

$v \leftarrow +\infty$

for each *a* **in** *game.ACTIONS*(*state*) **do**

$v2, a2 \leftarrow \text{MAX-VALUE}(\text{game}, \text{game.RESULT}(\text{state}, a))$

if $v2 < v$ **then**

$v, \text{move} \leftarrow v2, a$


return *v*, *move*

Compute and return min over
successors, all of which are MAX

Minimax Example


```
function MAX-VALUE(game, state) returns a (utility, move) pair
if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
v  $\leftarrow -\infty$ 
for each a in game.ACTIONS(state) do
    v2, a2  $\leftarrow$  MIN-VALUE(game, game.RESULT(state, a))
    if v2 > v then
        v, move  $\leftarrow$  v2, a
return v, move
```

```
function MIN-VALUE(game, state) returns a (utility, move) pair
if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
v  $\leftarrow +\infty$ 
for each a in game.ACTIONS(state) do
    v2, a2  $\leftarrow$  MAX-VALUE(game, game.RESULT(state, a))
    if v2 < v then
        v, move  $\leftarrow$  v2, a
return v, move
```



Non-Zero-Sum Games

- In non-zero-sum games, terminal states must record separate utilities for each player
- Each player maximizes own utility
- Different strategies may arise depending on the players' individual utility functions!
- May see cooperation, alliances, or team behavior if utilities are aligned




Minimax Efficiency

- Minimax is a form of DFS: Time complexity $O(b^m)$, space $O(bm)$
- Optimal if both players play perfectly, complete if game eventually ends
- Only games with sufficiently small game trees can be solved!
- Example: Chess has $b \approx 35$, $m \approx 80$
- We will need optimizations or approximations to speed up search
- Pruning the game tree, move ordering, depth limits, heuristic approaches


Pruning

- Suppose minimax is searching left to right
- We have finished searching branch A and computed its value (3)
- We start searching branch B and see value of one descendent
- Suppose its value is *smaller* than that of A (2)
- We know successors of root are all MIN nodes
- Regardless of other children, $v(B) \leq 2$
- Root MAX node will never consider B; skip (prune) and move on to C



Alpha-Beta Pruning

- Update two new values during search
- A node n is passed the highest (α) and lowest (β) values seen *along path* to n
- Skip remaining actions (prune) if:
 - MAX sees new value higher than β , since MIN parent will prefer β
 - MIN sees new value lower than α , since MAX parent will prefer α



Alpha-Beta Search

function ALPHA-BETA-SEARCH(*game*, *state*) **returns** an action

player \leftarrow *game*.TO-MOVE(*state*)

value, *move* \leftarrow MAX-VALUE(*game*, *state*, $-\infty$, $+\infty$)

return *move*

Assuming root is MAX

function MAX-VALUE(*game*, *state*, α , β) **returns** a (*utility*, *move*) pair

if *game*.IS-TERMINAL(*state*) **then return** *game*.UTILITY(*state*, *player*), *null*

v $\leftarrow -\infty$

for each *a* **in** *game*.ACTIONS(*state*) **do**

v2, *a2* \leftarrow MIN-VALUE(*game*, *game*.RESULT(*state*, *a*), α , β)

if *v2* > *v* **then**

v, *move* \leftarrow *v2*, *a*

$\alpha \leftarrow \text{MAX}(\alpha, v)$

if *v* $\geq \beta$ **then return** *v*, *move*

return *v*, *move*

MAX updates α ; if node value is $\geq \beta$,
prune remaining successors

function MIN-VALUE(*game*, *state*, α , β) **returns** a (*utility*, *move*) pair

if *game*.IS-TERMINAL(*state*) **then return** *game*.UTILITY(*state*, *player*), *null*

v $\leftarrow +\infty$

for each *a* **in** *game*.ACTIONS(*state*) **do**

v2, *a2* \leftarrow MAX-VALUE(*game*, *game*.RESULT(*state*, *a*), α , β)

if *v2* < *v* **then**

v, *move* \leftarrow *v2*, *a*


$\beta \leftarrow \text{MIN}(\beta, v)$

if *v* $\leq \alpha$ **then return** *v*, *move*

return *v*, *move*


MIN updates β ; if node value is $\leq \alpha$,
prune remaining successors

Alpha-Beta Example



Alpha-Beta Properties

- Pruning still results in correct minimax value at root
- Intermediate (children) node values might be wrong!!
- If pruning, game tree values cannot be reused
- Will have to rerun minimax after each move
- In practice, we can store just the states/values that we know to be correct in a *transposition table* in case they come up again
 - Especially effective when there are multiple paths to a state



Move Ordering

- Good *move ordering* improves effectiveness of pruning
- Alpha-beta with random ordering is roughly $\sim O(b^{0.75m})$
- “Perfect ordering” gets us to $O(b^{0.5m})$, doubling solvable depth
- Usually based on domain or expert knowledge
 - Simple chess ordering function: captures, threats, forward moves, backward moves
- Iterative deepening: Since upper game tree layers will be expanded multiple times, can also keep track of their values in order to optimally order them in future iterations
- Search of repeated states can be avoided by caching them in a transposition table

Summary

- Adversarial search can be used to solve multi-agent problems
- Minimax is well suited for two-player, zero-sum, and turn-taking games
- Utility values computed at leaf nodes, backed up to non-end states
- Alpha-beta pruning can make more search more efficient by avoiding useless portions of the game tree
- Effective move ordering and caching can further improve search efficiency