

# COMS W4701: Artificial Intelligence

## Lecture 8: Heuristic Game Playing

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences


# Today

---

- Depth limits and evaluation functions
- Stochastic games
- Monte Carlo tree search

# Imperfect Decisions

- Problem: Most game trees still too big
- $\alpha$ - $\beta$  can help but still need to find terminal nodes
- Heuristic: Treat non-terminal nodes as terminals!
- **Evaluation function** returns an *estimate* of this “terminal” state’s utility
- **Cutoff test** decides when to do this
- No more guarantee of optimality




# Evaluation Functions

---

- Evaluation functions are *estimates* of a state's utility
- Evaluation functions...
  - Should be equal to utility values for terminal nodes
  - Should be reflective of minimax values for nonterminal nodes
  - Should be efficient to compute and based on game knowledge/rules
- One common eval function: weighted linear sum of game **features**
- Features represent categories or equivalence classes of states

# Example: Chess



$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- Features may be derived from expert knowledge of common *categories* of states
- E.g., one feature for each type of piece, attack formations, king safety positions, etc.
- Weights correspond to *material values* of each feature
- Linear weighting assumes features are independent of each other

# Depth Limits

---


- What cutoff test to use?
- Simple approach: Use a fixed depth limit, or use iterative deepening
- Problem: Cutting off search at unstable positions can hide valuable info
- **Horizon effect:** Agent may favor moves that push danger “over the horizon”, appears to have been mitigated but actually delayed
- **Quiescence search:** Evaluate states based on their stability (“quiescence”), e.g., chess moves that do not lead to imminent captures

# Stochastic Games


- Many games contain elements of *stochasticity*
- Opponents playing suboptimally (e.g., inexperienced) or randomly
- Explicit random elements (e.g., dice rolling)
- Instead of worst-case scenarios, we consider **expected outcomes**
- Instead of minimax, we compute **expectiminimax values**

EXPECTIMINIMAX( $s$ ) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \\ \sum_r P(r) \text{EXPECTIMINIMAX}(\text{RESULT}(s, r)) & \text{if } \text{PLAYER}(s) = \text{CHANCE} \end{cases}$$




# Example: Backgammon





# Heuristics for Expectiminimax

- With deterministic games,  $\alpha$ - $\beta$  can cleverly prune the game tree and ignore outcomes that cannot occur
- With stochastic games, *any* outcomes may occur
- Even the unlikely outcomes must be examined
- We may be able to prune if we can *bound* utility values
- We can also *forward-prune* by ignoring bad looking moves a priori (as in beam search)



All utilities  $\leq 10$

# Monte Carlo Tree Search

---

- Minimax and its variants are **type A** strategies: search wide but shallow
- For games like Go with large branching factors, **type B** strategies like MCTS work better: search deep but narrow
- *Traverse* and *expand* new nodes in promising portions of the game tree
- For a partial tree, *simulate* the game from a leaf node to a terminal
- *Backpropagate* the result up to all nodes along traversed path
- **MCTS**: Repeat the above many times and choose the most traversed move



# Selection


- We traverse the tree until arriving at a terminal node or one with successors not yet in the tree
- A **selection policy** balances *exploitation* (highest values) with *exploration* (less visited nodes)
- **UCT** method for traversal: From each node, move to successor maximizing the following:

$$UCT(n) = \underbrace{\frac{w(n)}{N(n)}}_{\text{“exploitation”}} + \alpha \underbrace{\sqrt{\frac{\ln N(\text{parent}(n))}{N(n)}}}_{\text{“exploration”}}$$

“exploitation”

“exploration”

$\alpha$  = exploration parameter



$w(n)$  = number of wins from  $n$

$N(n)$  = number of rollouts from  $n$

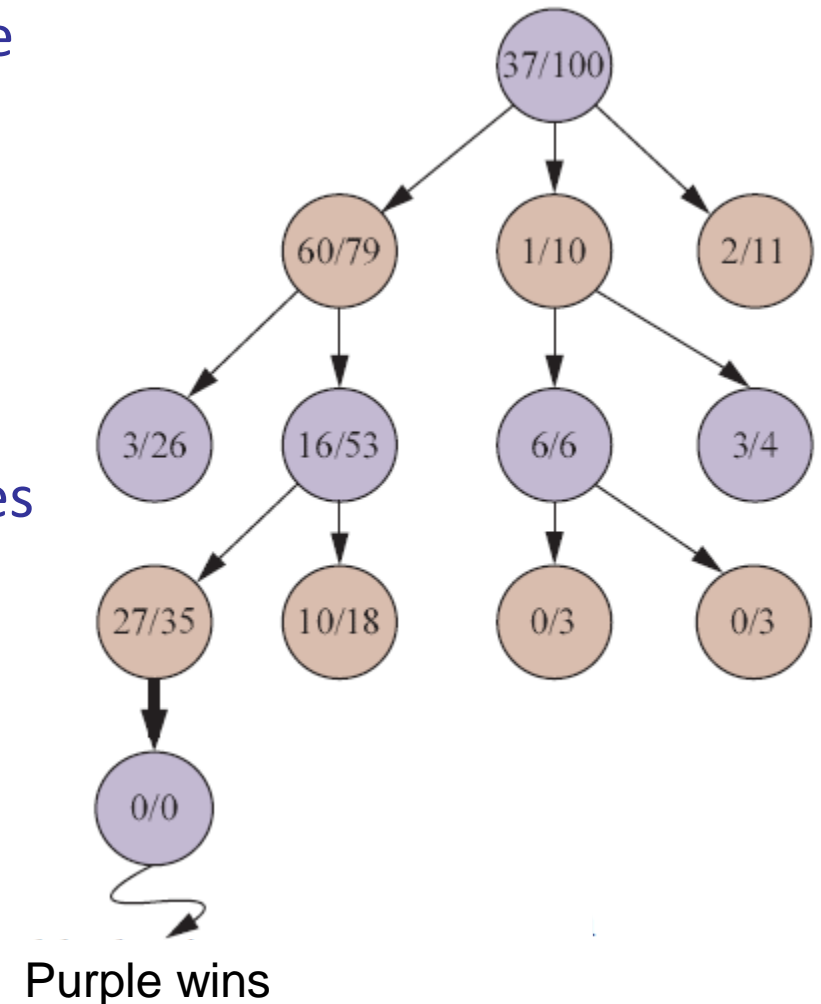
# Upper Confidence Bound

---

- The *exploration* or *uncertainty* term of the UCT value is  $\alpha \sqrt{\frac{\ln N(\text{parent}(n))}{N(n)}}$
- $\alpha \geq 0$ : Tunable parameter controlling weight of this term
- $\sqrt{\ln(\text{parent}(N))}$ : Proportional to number of rollouts of parent node, can just treat as a constant since same value for all other children nodes
- $N(n)$  is the number of rollouts containing  $n$ , or *samples* for estimating the value of  $n$
- $1/\sqrt{N(n)}$  is proportional to the standard deviation of this estimate
- Each rollout through  $n$  increments  $N(n)$  and *decreases* its uncertainty

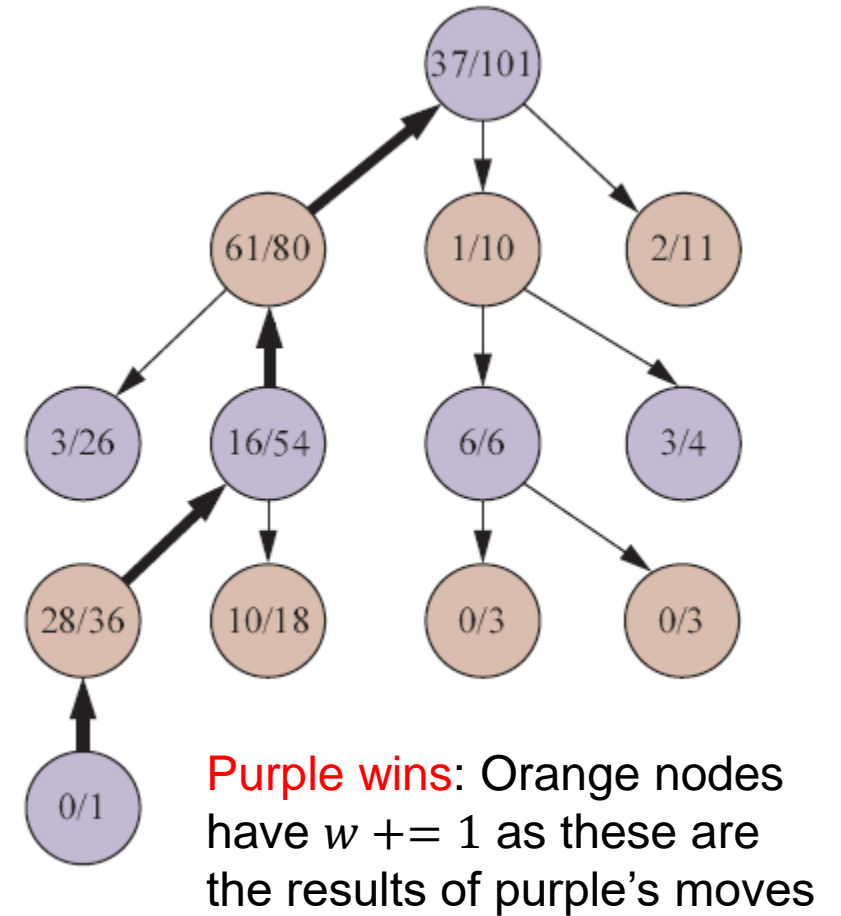
# Expansion and Simulation

- Expansion: Add a child of the selected node to the tree
- **Playout policy** simulates a *rollout* from the new child
- Basic example: Select moves at random
- Better policies guide moves toward good or clever ones
- May incorporate game-specific heuristics
- May be obtained from deep learning (e.g. AlphaGo)
- Record the game result, but *not* the game states seen

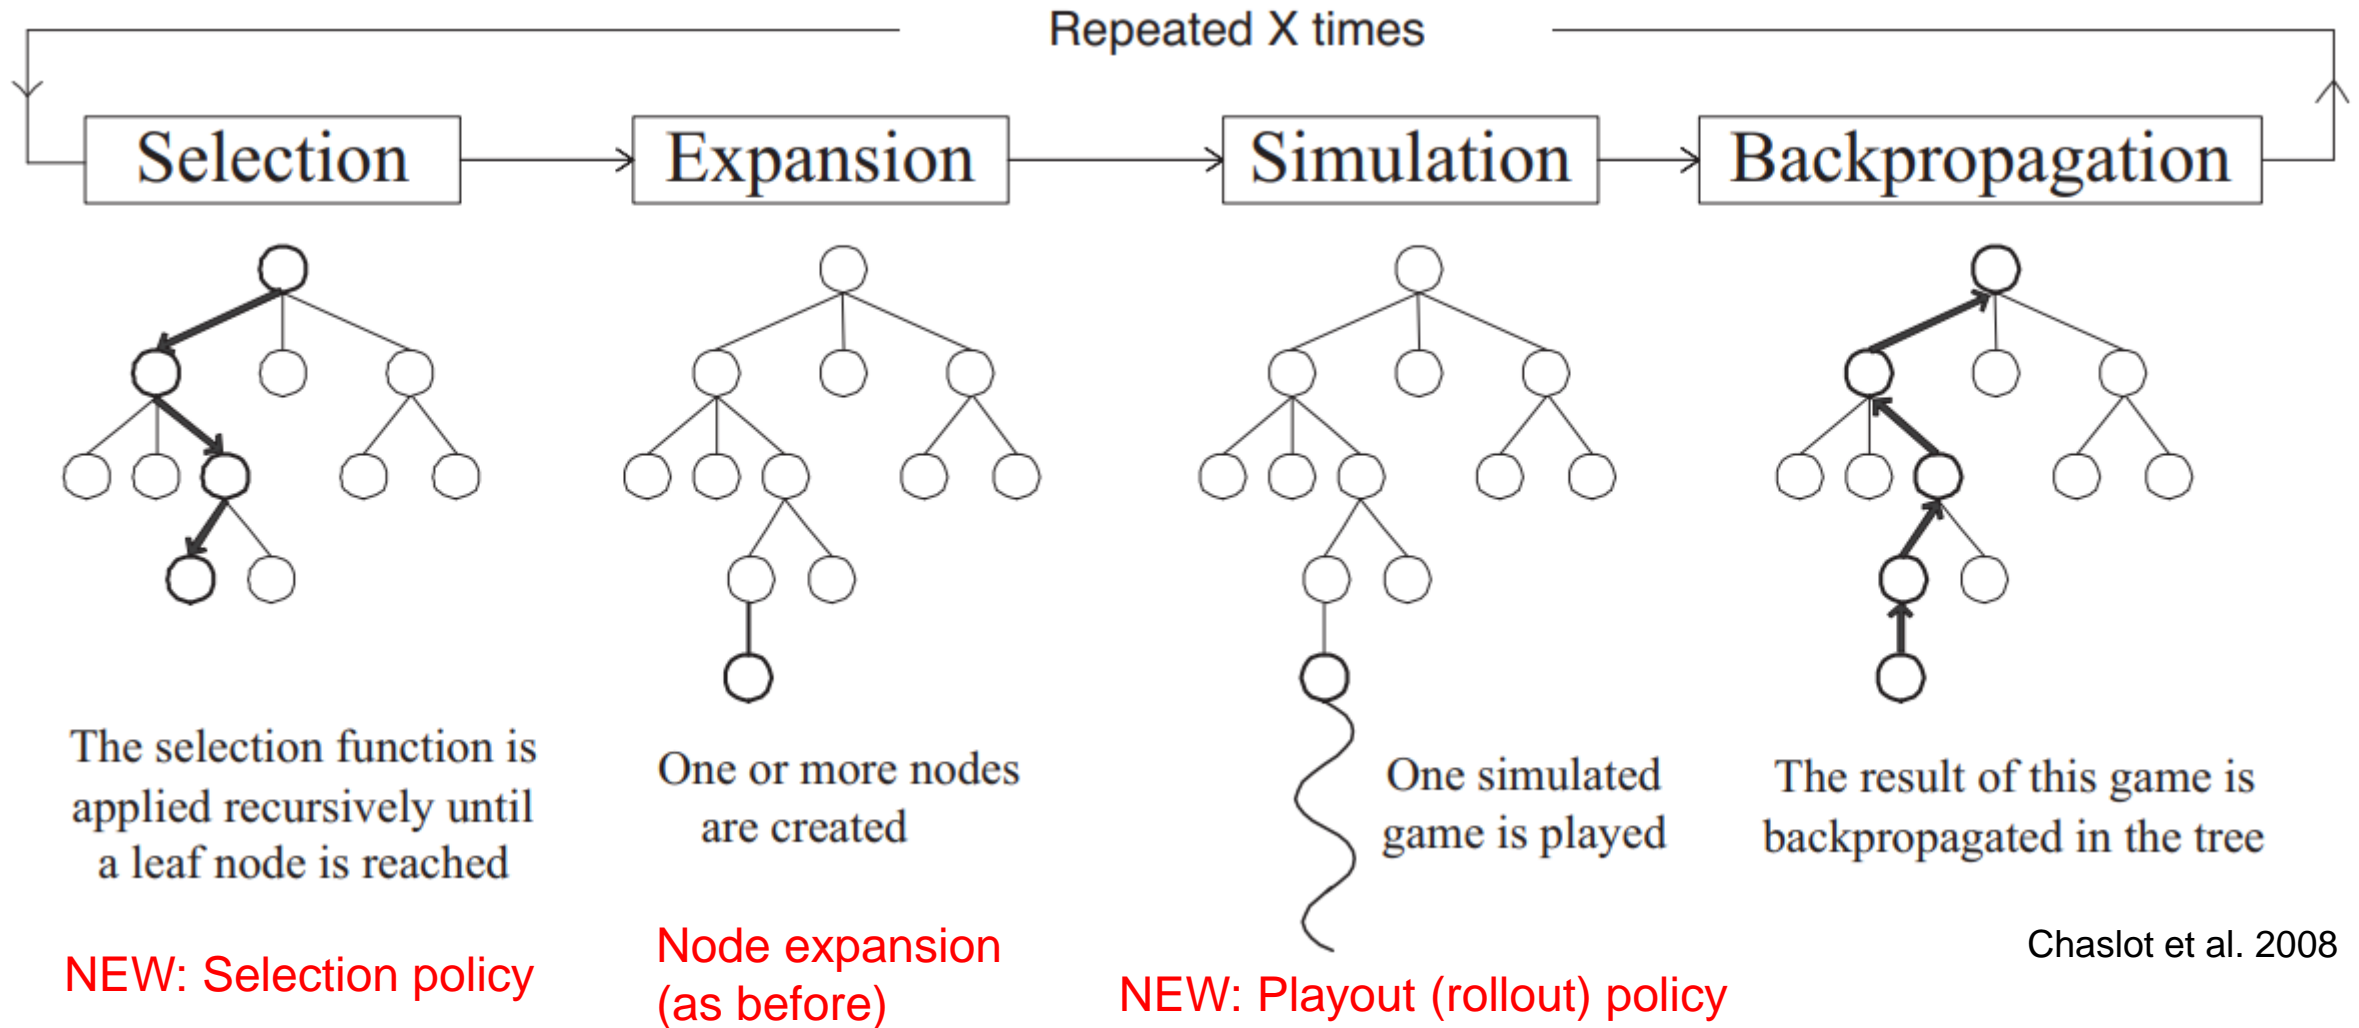


# Backpropagation

- Once simulation finishes, pass information back up to *all* nodes along path back to root
- Increment  $N$  value of all nodes on path
- If tracking wins, increment  $w(n)$  of nodes whose *parent* is the winning player
- After sufficient rounds of MCTS, the root chooses the most frequent action (highest  $N$ ) as the move to play
- More robust and less variability than the move with the highest value (if different)



# Monte Carlo Tree Search





# MCTS vs Alpha-Beta

---

- MCTS simulations are linear in game depth
- For a game with branching factor of 32 and depth of 100, alpha-beta search down to 12 ply deep is equivalent to  $10^7$  MCTS simulations
- MCTS tends to do better when branching factor is high
- Also less sensitive than  $\alpha$ - $\beta$  to inaccurate eval functions
- Also good for brand new games with no predefined eval functions at all!
- Stochastic nature of MCTS: no guarantee of exploring all good moves

# Summary

---

- Practical adversarial search often means exploring portions of game tree
- Alpha-beta search can be cut off at finite depth limits by applying evaluation functions to non-terminal states
- Games with stochasticity can be handled with expectiminimax search
- Games with high branching factors can be searched deep, not wide
- MCTS builds a partial tree over many iterations, simulating many playouts to the end of the game