

COMS W4701: Artificial Intelligence

Lecture 10: Dynamic Programming

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences

Today


- Time-limited values
- Iterative policy evaluation
- Value iteration

MDPs and Bellman Equations

- **MDPs:** Stochastic, sequential decision problems
- **Policy** $\pi: S \rightarrow A$, assignment of action to each state
- **Value** $V^\pi: S \rightarrow \mathbb{R}$, expected utility starting at each state and following π
- **Bellman optimality equations:**

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$



Time-Limited Values


- Math problem: State values are nonlinear functions of other state values
- We can try traversing the tree using search, but may be exponentially large
- Idea from depth-limited search: Treat non-terminal states as terminal
- If the time horizon is 0, *every* state is effectively “terminal”
- **Time-limited values** $V_t(s)$: Expected utility of t decisions starting from s
- If we have $V_t(s)$, we can perform a tree value backup to compute $V_{t+1}(s)$

Time-Limited Values

- Given a policy π , we can *iterate* over time-limited values to solve for V^π
- Alternative method to fully solving a system of linear equations
- Base case: Time horizon $i = 0$. No more rewards, so $V_0(s) = 0 \forall s$
- $i = 1$: $V_1^\pi(s) = \sum_{s'} T(s, \pi(s), s') R(s, \pi(s), s')$
- $i = 2$: $V_2^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_1^\pi(s')]$
- $i = 3$: $V_3^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_2^\pi(s')]$
- ...

Iterative Policy Evaluation

- Using V_i to iteratively compute V_{i+1} is an example of *dynamic programming*
- Initialize $V_0^\pi(s) \leftarrow 0$ for all states s
- **Loop** from $i = 0$:
 - **Initialize** temporary array V_{i+1}
 - **For** each state $s \in S$:
$$V_{i+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_i^\pi(s')]$$
 - **Copy** V_{i+1} into V_i
- **Until** $\max_s |V_{i+1}^\pi(s) - V_i^\pi(s)| < \epsilon$ (small threshold)



Example: Mini-Gridworld

$$V_{i+1}^{\pi}(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_i^{\pi}(s')]$$

+3	-2	+1
<i>A</i>	<i>B</i>	<i>C</i>



- Suppose we are given the policy $\pi(s) = L \ \forall s$
- Suppose we use the discount factor $\gamma = 0.5$
- We can iteratively perform three value updates, one for each state:

$$V_{i+1}^{\pi}(A) = 0.8(3 + 0.5V_i^{\pi}(A)) + 0.2(-2 + 0.5V_i^{\pi}(B))$$

$$V_{i+1}^{\pi}(B) = 0.8(3 + 0.5V_i^{\pi}(A)) + 0.2(1 + 0.5V_i^{\pi}(C))$$

$$V_{i+1}^{\pi}(C) = 0.8(-2 + 0.5V_i^{\pi}(B)) + 0.2(1 + 0.5V_i^{\pi}(C))$$

Example: Mini-Gridworld




Value Iteration

- Now suppose we want to find V^* rather than evaluate a fixed policy
- Additional step: Find the value of the *best* action in each iteration
- Initialize: $V_0(s) \leftarrow 0$ for all states s
- **Loop** from $i = 0$:
 - **Initialize** temporary array V_{i+1}
 - **For** each state $s \in S$:

Bellman update

$$V_{i+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i(s')]$$
 - **Copy** V_{i+1} into V_i
- **Until** $\max_s |V_{i+1}(s) - V_i(s)| < \epsilon$ (small threshold)



Example: Mini-Gridworld

$$V_{i+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i(s')]$$

+3	-2	+1
<i>A</i>	<i>B</i>	<i>C</i>



- States, actions, rewards as shown; $\gamma = 0.5$
- Transitions: $\text{Pr}(\text{intended direction}) = 0.8$, $\text{Pr}(\text{opposite direction}) = 0.2$

$$V_{i+1}(A) = \max \begin{bmatrix} 0.8(3 + 0.5V_i(A)) + 0.2(-2 + 0.5V_i(B)), & L \text{ action} \\ 0.8(-2 + 0.5V_i(B)) + 0.2(3 + 0.5V_i(A)) & R \text{ action} \end{bmatrix}$$

$$V_{i+1}(B) = \max \begin{bmatrix} 0.8(3 + 0.5V_i(A)) + 0.2(1 + 0.5V_i(C)), & L \text{ action} \\ 0.8(1 + 0.5V_i(C)) + 0.2(3 + 0.5V_i(A)) & R \text{ action} \end{bmatrix}$$

$$V_{i+1}(C) = \max \begin{bmatrix} 0.8(-2 + 0.5V_i(B)) + 0.2(1 + 0.5V_i(C)), & L \text{ action} \\ 0.8(1 + 0.5V_i(C)) + 0.2(-2 + 0.5V_i(B)) & R \text{ action} \end{bmatrix}$$

Example: Mini-Gridworld



Convergence of Value Iteration

- The Bellman update is a **contraction mapping**
- Time-limited value are always guaranteed to move toward V^*
- Fact 1: Bellman update does not change optimal values V^* (*fixed point*)

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$


- Define the *max norm* (error) $\|V_i - V^*\| = \max_s |V_i(s) - V^*(s)|$
- Fact 2: Each update shrinks max error in V by factor of γ : $\|V_{i+1} - V^*\| \leq \gamma \|V_i - V^*\|$
- Errors decrease (and values converge) exponentially fast!

Policy Extraction

- How do we back out π^* after computing V^* ?
- (Recursive) definition of π^* from Bellman equation:

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- Everything on the RHS is now known!
- For each state, iterate through all actions
- Optimal policy assigns action with the highest utility



Algorithm Complexity

$$V_{i+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i(s')]$$

- Each sweep of value iteration involves, at each state, an expectation over all successor states and a maximization over all actions
- Same operation for policy extraction, so each sweep takes $O(|S|^2|A|)$ time
- The *number* of sweeps depends on discount factor γ and error threshold ϵ
- Since min error reduction is γ per iteration, a forward-looking agent (larger γ) requires more sweeps before convergence

Summary

- Dynamic programming solves MDPs using time-limited quantities derived from decision making in finite time horizons
- Bellman updates push values and policies toward the optimal solution
- Policy evaluation: Iteratively update and solve for values of fixed policy
- Value iteration: Compute optimal values for all states by iteratively finding best actions and their values at each iteration