

COMS W4701: Artificial Intelligence

Lecture 24: Linear Regression and Classification

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences

Today

- Linear regression
- Linear classification
- Perceptron learning rule
- Logistic regression
- Gradient descent



Linear Models

- Consider a supervised learning task in which input data are *feature vectors* $\mathbf{x}_i = (\mathbf{x}_{i1}, \dots, \mathbf{x}_{ip})$ and outputs are *scalars* y_i

- A **linear model** is described by a **weight vector** \mathbf{w} :

$$f_{\mathbf{w}}(\mathbf{x}_i) = w_0 + \sum_{j=1}^p w_j x_{ij} = \mathbf{w} \cdot (1, \mathbf{x}_i)$$

- Data sets are generally noisy, and it is generally impossible to recover a hypothesis \hat{f} s.t. $\hat{f}(\mathbf{x}_i) = y_i$ for all i
- Linear regression** tries to find a hypothesis so as to *minimize the loss* on the data set





Loss Functions

- A **loss function** $L(y, \hat{y})$ measures the closeness between a predicted value \hat{y} and the actual value y
 - Absolute (L_1) loss: $L(y, \hat{y}) = |y - \hat{y}|$
 - Squared (L_2) loss: $L(y, \hat{y}) = (y - \hat{y})^2$
 - **0-1** loss: $L(y, \hat{y}) = 1$ if $y \neq \hat{y}$, and 0 otherwise
- The *empirical loss* can be computed over a data set can be computed by taking the sum or the mean of one of the functions above
- L_1 and L_2 losses are useful for real values, 0-1 loss for categorical values

Linear Regression

- Given a data set $((\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n))$, we can stack the feature vectors into a $n \times p$ **data matrix** \mathbf{X} and the output values y_i into a n -length **output vector** \mathbf{y}
- Given a p -length **weight vector** \mathbf{w} , we can compute the *predictions* of all \mathbf{x}_i as $\hat{\mathbf{y}} = \hat{f}_{\mathbf{w}}(\mathbf{X}) = \mathbf{X}\mathbf{w}$
- Linear regression minimizes the *total squared-error loss* $\sum_i (y_i - \hat{y}_i)^2$, and has a closed-form solution!

$$\mathbf{w}^* = \min_{\mathbf{w}} \|\mathbf{y} - \hat{\mathbf{y}}\|^2 = \min_{\mathbf{w}} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2 = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

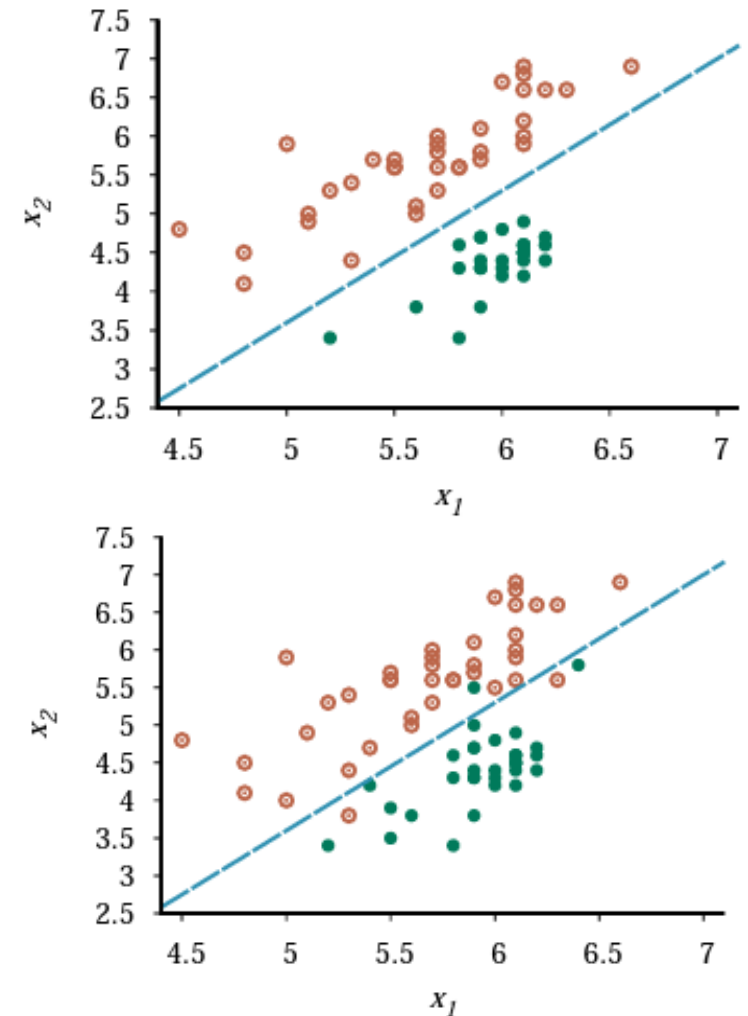


Linear Classification

- Now consider a **classification** problem, in which output values y are categorical, e.g. $\{0,1\}$
- One way of classifying data is to *separate* them in feature space, e.g. using linear functions
- For binary classification, we can define an **activation function** that composes with the linear function:


$$h_{\mathbf{w}}(\mathbf{x}) = \begin{cases} 1, & \text{if } f_{\mathbf{w}}(\mathbf{x}) \geq 0 \\ 0, & \text{if } f_{\mathbf{w}}(\mathbf{x}) < 0 \end{cases}$$

- The hyperplane $h_{\mathbf{w}}(\mathbf{x}) = 0$ forms the **decision boundary**



Example: Linear Classifier

- Given an email input, $\mathbf{x} = (\text{Contains? ("free")}, \text{Contains? ("money")})$
- Want to predict either $y = 1$ (spam) or $y = 0$ (ham)
- Suppose our model (weight vector) is $\mathbf{w} = (-3, 4, 2)$
- Separating hyperplane is $f_{\mathbf{w}}(\mathbf{x}) = -3 + 4x_1 + 2x_2 = 0$
- “Free money”: $h_{\mathbf{w}}(\mathbf{x}) = h_{\mathbf{w}}((1,1)) = 1 \rightarrow \text{spam}$
- “Free food”: $h_{\mathbf{w}}(\mathbf{x}) = h_{\mathbf{w}}((1,0)) = 1 \rightarrow \text{spam}$
- “No money”: $h_{\mathbf{w}}(\mathbf{x}) = h_{\mathbf{w}}((0,1)) = 0 \rightarrow \text{ham}$
- “how are you”: $h_{\mathbf{w}}(\mathbf{x}) = h_{\mathbf{w}}((0,0)) = 0 \rightarrow \text{ham}$



Perceptron Learning Rule

- No closed-form solution! Activation function makes the problem nonlinear
- *Iterative* method to learn \mathbf{w} : Classify each training data instance, and update \mathbf{w} if prediction is incorrect

Initialize weights \mathbf{w} (e.g., all 0) and learning rate α


While not converged:

For each training instance \mathbf{x}_i :

 Predict class y using current weights \mathbf{w}


If incorrect ($y \neq y_i^*$), update weights:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha(y_i^* - y)(1, \mathbf{x}_i)$$




Example: Perceptron Learning Rule

- Suppose we have the following training data:
- $((0,1), \text{spam}), ((1,0), \text{ham}), ((0,0), \text{ham})$
- We currently have $\mathbf{w} = (-3, 4, 2)$ and we use $\alpha = 1$
- $h_{\mathbf{w}}(\mathbf{x}_1) = 0 \rightarrow \text{ham} \rightarrow \mathbf{w} = \mathbf{w} + (1, 0, 1) = (-2, 4, 3)$
- $h_{\mathbf{w}}(\mathbf{x}_2) = 1 \rightarrow \text{spam} \rightarrow \mathbf{w} = \mathbf{w} - (1, 1, 0) = (-3, 3, 3)$
- $h_{\mathbf{w}}(\mathbf{x}_3) = 0 \rightarrow \text{ham}$; correct, no change
- $h_{\mathbf{w}}(\mathbf{x}_1) = 1 \rightarrow \text{spam}$; correct, no change
- $h_{\mathbf{w}}(\mathbf{x}_2) = 1 \rightarrow \text{spam} \rightarrow \mathbf{w} = \mathbf{w} - (1, 1, 0) = (-4, 2, 3)$
- ... until all data classified correctly




Perceptron Convergence


- For linearly separable data, the perceptron is *guaranteed* to find a linear separator (though convergence may not be monotonic)
- For nonseparable data, perceptron will converge to a minimum-error solution if α is decayed as $O\left(\frac{1}{t}\right)$



Separable data




Nonseparable data



Nonseparable data
with decaying α



Logistic Model

- Hard threshold causes training to be discontinuous and unstable
- Classifier is not informed by the distance between data and boundary
- It is possible for the perceptron to learn *barely separating* boundaries
- Idea: Replace the *hard* threshold with a *soft* threshold
- Use a **logistic** (sigmoid) function for h :
$$h_{\mathbf{w}}(\mathbf{x}) = \sigma(f_{\mathbf{w}}(\mathbf{x})) = \frac{1}{1 + \exp(-f_{\mathbf{w}}(\mathbf{x}))}$$
- Outputs $[0,1]$ can be interpreted as probabilities!
- Large values of $f_{\mathbf{w}}(\mathbf{x})$ push $h_{\mathbf{w}}(\mathbf{x})$ closer to 1
- On the decision boundary $f_{\mathbf{w}}(\mathbf{x}) = 0$, we have $h_{\mathbf{w}}(\mathbf{x}) = 0.5$



Example: Logistic Model

- Weights: $\mathbf{w} = (-10.65, .0055)$
- Logistic model: $h_{\mathbf{w}}(x) = \frac{1}{1+\exp(-(-10.65+.0055x))}$
- We can also say $\Pr(y = 1|x) = h_{\mathbf{w}}(x)$
- $\Pr(\text{default}|\text{balance} = 1000) = 0.00576 \rightarrow \text{predict no default}$
- $\Pr(\text{default}|\text{balance} = 2000) = 0.586 \rightarrow \text{predict default}$
- Note that $\log \frac{\Pr(Y=1|x)}{\Pr(Y=0|x)} = \log \frac{h_{\mathbf{w}}(x)}{1-h_{\mathbf{w}}(x)} = w_0 + wx$
- The **log odds** is exactly the linear function $f_{\mathbf{w}}(x)$



Logistic Regression

- Like the hard threshold, no closed form solution for learning weights
- Unlike the hard threshold, there is a well-defined iterative approach based on *maximizing the likelihood* of our data
- For (\mathbf{x}_i, y_i) , maximize the probability $\Pr(y_i|\mathbf{x}_i) = h_{\mathbf{w}}(\mathbf{x}_i)$ if $y_i = 1$, and alternatively $1 - h_{\mathbf{w}}(\mathbf{x}_i)$ if $y_i = 0$
- In other words, maximize the likelihood

$$\Pr(y_i|\mathbf{x}_i) = h_{\mathbf{w}}(\mathbf{x}_i)^{y_i} (1 - h_{\mathbf{w}}(\mathbf{x}_i))^{1-y_i}$$

Log Likelihood

- We will write the *negative log likelihood* to obtain easier derivatives:

$$-\log h_{\mathbf{w}}(\mathbf{x}_i)^{y_i} (1 - h_{\mathbf{w}}(\mathbf{x}_i))^{1-y_i} = -(y_i \log h_{\mathbf{w}}(\mathbf{x}_i) + (1 - y_i) \log(1 - h_{\mathbf{w}}(\mathbf{x}_i)))$$

- We will *minimize* the average negative log likelihood over all data:

$$L(\mathbf{w}) = -\frac{1}{n} \sum_{i=1}^n (y_i \log h_{\mathbf{w}}(\mathbf{x}_i) + (1 - y_i) \log(1 - h_{\mathbf{w}}(\mathbf{x}_i)))$$

- We will need the partial derivatives of L wrt each of its weights w_i :



$$\frac{\partial L}{\partial w_0} = \frac{1}{n} \sum_{i=1}^n (h_{\mathbf{w}}(\mathbf{x}_i) - y_i) \quad \frac{\partial L}{\partial w_i} = \frac{1}{n} \sum_{i=1}^n (h_{\mathbf{w}}(\mathbf{x}_i) - y_i) \mathbf{x}_i$$

Gradient Descent

- We are *searching* for a weight vector \mathbf{w} that minimizes L
- The **gradient** of L is a *vector* of partial derivatives wrt all w_j

$$\frac{\partial L}{\partial \mathbf{w}} = \left(\frac{\partial L}{\partial w_0}, \frac{\partial L}{\partial w_1}, \dots, \frac{\partial L}{\partial w_p} \right)$$

- Indicates magnitude and *direction* of largest increase in L
- The value of L *decreases* fastest along direction of $-\frac{\partial L}{\partial \mathbf{w}}$
- **Gradient descent:** Initialize the configuration \mathbf{w} , and repeatedly update it as $\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial L}{\partial \mathbf{w}}$ until convergence



Stochastic Gradient Descent



- $\frac{\partial L}{\partial \mathbf{w}}$ is computed by summing classification “errors” over all data:

$$\frac{\partial L}{\partial \mathbf{w}} = \frac{1}{n} \sum_{i=1}^n (h_{\mathbf{w}}(\mathbf{x}_i) - y_i) \begin{bmatrix} 1 \\ \mathbf{x}_i \end{bmatrix}$$

- **Batch gradient descent** thus computes a single weight vector update per pass through the data set (**epoch**), opposite behavior of perceptron (update for each *data instance*)
- **Stochastic gradient descent** *approximates* $\frac{\partial L}{\partial \mathbf{w}}$ by selecting or sampling a **minibatch** of size m from the data set to compute gradient and update weights
- One epoch will thus see $\frac{n}{m}$ weight updates, allowing for faster convergence

SGD for Logistic Regression

- Given data set $\mathbf{d} = ((\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n))$, minibatch size m , learning rate α
- Initialize weights \mathbf{w} , e.g. randomly
- **Until** \mathbf{w} has converged:
 - Sample minibatch MB from \mathbf{d}
 - $grad = \mathbf{0}$
 - **For** $(\mathbf{x}, y) \in MB$:
 - $grad = grad - (y - h_{\mathbf{w}}(\mathbf{x})) \begin{bmatrix} 1 \\ \mathbf{x}_i \end{bmatrix}$
 - $\mathbf{w} = \mathbf{w} - \frac{\alpha}{m} \cdot grad$



Multinomial Logistic Regression

- If we have K classes, then we can use the **softmax** function
- Learn a weight vector \mathbf{w}_k for each class k , which can be used to compute the probability of each class given an input \mathbf{x}

$$\Pr(Y = k|\mathbf{x}) = \frac{\exp(\mathbf{w}_k \cdot (1, \mathbf{x}))}{\sum_{i=1}^K \exp(\mathbf{w}_i \cdot (1, \mathbf{x}))}$$

- Classifier can output the most likely class or the distribution itself
- The points where $\Pr(Y = k_1|\mathbf{x}) = \Pr(Y = k_2|\mathbf{x})$ form a linear decision boundary between the two classes k_1 and k_2

Summary

- Linear models compute linear combinations of feature vectors
- Closed-form solution for weight vector in the regression problem
- Linear classifiers find a decision boundary in feature space
- Can apply an activation function to the linear function output
- The perceptron learning rule can be used for hard thresholds
- Logistic regression applies a *soft* threshold to the weighted feature
- Can be solved by gradient descent of likelihood function